Compositional design of asynchronous circuits from behavioural concepts

Jonathan Beaumont, Andrey Mokhov, Danil Sokolov, Alex Yakovlev
{j.r.beaumont, andrey.mokhov, danil.sokolov, alex.yakovlev}@ncl.ac.uk
School of Electrical and Electronic Engineering, Newcastle University, UK

Abstract—**Asynchronous circuits can be useful in many applications, however, they are yet to be widely used in industry. The main reason for this is a steep learning curve for concurrency models, such Signal Transition Graphs, that are developed by the academic community for specification and synthesis of asynchronous circuits. In this paper we introduce a compositional design flow for asynchronous circuits using *concepts* – a set of formalised descriptions for system requirements. Our aim is to simplify the process of capturing system requirements in the form of a formal specification, and promote the concepts as a means for design reuse. The proposed design flow is applied to the development of an asynchronous buck converter.**

1. Introduction

Asynchronous circuits are event-driven, i.e. they react to changes in a system at the rate they occur [1]. This makes them particularly useful for on-chip power management, where the ability to quickly respond to dynamically changing loads across the chip is essential for reliable operation and efficiency [2]. A power management system relies on analogue circuitry for power regulation and conversion whose behaviour is characterised by many operating modes and complexity of their interplay. Capturing all these aspects of system behaviour in a consistent specification becomes the major design challenge [3].

Signal Transition Graphs (STGs) are commonly used for the specification of asynchronous control circuits as they are compatible with multiple synthesis tools, such as PETRIFY [4] and MPSAT [5]. These tools take an STG specification of a complete controller and produce a speed-independent circuit implementation [6]. Such a monolithic approach to designing asynchronous circuits has poor scalability: as the system grows in complexity its monolithic specification becomes challenging to comprehend and debug. The STG models of components cannot be reused when designing other specifications, and thus each new design must be built from the ground up. This further adds to the design time, hence making asynchronous circuits undesirable for use in industry.

To address this issue, we propose a new method of asynchronous circuit design. The method splits a specification into several parts corresponding to operational modes of the circuit (*scenarios*). The features, constraints and requirements of each scenario (*concepts*), are described in a formal

notation, which we implemented as a domain specific language embedded in Haskell [7]. Concepts can be composed and one concept can be made up of multiple smaller concepts, thus supporting the design reuse at the level of system specification. Scenarios of reconfigurable systems [8] can also be parameterised by run-time parameters (e.g., available energy budget) or design-time ones (e.g., the number of processing cores in a Network-on-Chip network), therefore concepts should also support parameterisation.

A set of concepts describing the operation of a scenario is then passed into a translation algorithm that automatically converts it into an equivalent STG, which satisfies all given concepts and can be model-checked using standard tools [9]. When all scenarios have been translated to STGs and verified, they can be combined to produce a complete specification. This step will also be automated, and will offer *templates* for common scenario ordering requirements, such as mode switching sequences and start-up scenarios.

Designing a controller for an analogue circuit using this method can be beneficial. Any of the partial knowledge we have about any casual relationships between events in the environment can be naturally modelled as concepts. When composed with other concepts describing these relationships and concepts describing the control which reacts to the environment, a model will be produced which shows how the environment and the control system interacts.

The idea is that our approach should reduce the complexity of designing asynchronous circuits, so the number of errors should be reduced, and easier to find and correct. This will in turn reduce the design time, and make asynchronous circuits more desirable and be used to make devices.

The presented approach is automated in the open-source WORKCRAFT framework [9]. This parses concepts, uses them to produce scenario STGs and performs parallel composition [15] on these and creates full model STGs. Full models can then be synthesised using WORKCRAFT. In this paper, we use real life industrial example of a buck converter [3] to show our design flow, and test this design approach.

2. Motivating Example

Signal transition graphs are commonly used for the specification of asynchronous circuits, and a monolithic design approach is used when designing with STGs. This approach

can become problematic however as designs become larger. In this section we will discuss this approach in detail and it's disadvantages. Following this, we will explore a way of viewing STGs which has helped us develop concepts.

2.1. Monolithic design approach for asynchronous circuits

The example we will use throughout this paper is that of a simple buck converter. A detailed description of this can be found in Section 6, but for this section it serves as an example of how the monolithic design approach is carried out. Figure 1 contains the final STG for a simple buck converter, and we will discuss how this STG is attained using the monolithic approach.
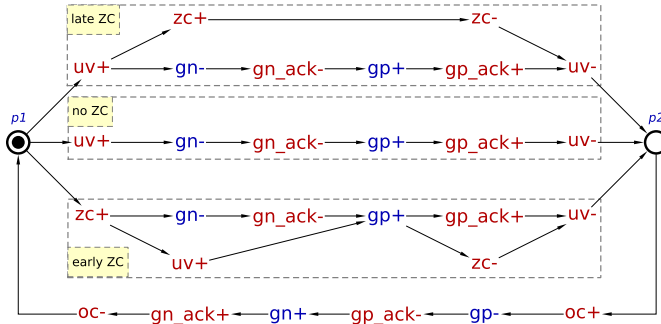

Figure 1: STG of a simple buck converter.

This STG is quite large, and has several interconnected branches. The monolithic approach starts with a blank page, and a designer would have to manually insert signals and the connections between these according to the description of operation of the system. With there being several branches, the designer could choose to design each branch separately, and manually connect these depending on the similarities between these branches.

In the event that a designer needs to add or remove signals or correct a fault with this design, editing this design can become difficult, due to the complexity and size. This can lead to further faults, and several iterations design and tests until the new feature is added and deemed to be working correctly, or the designer may even have to start from a blank page, unable to reuse any of the previous design. The larger and more complex an STG is, and the more signals there are, the more difficult it is to comprehend, debug and edit. This can slow the design process of a circuit, which is undesirable in an industry where the time for a device to move from conception to market is becoming critical, and ever shorter [3].

2.2. A new way of viewing STGs

In an attempt to streamline the design process, we aimed to find a way to create STGs which allows editing at various stages of the design. This way, if something needs to be changed, this can be done to a small part, but we can reuse

everything that works correctly. These can then be composed to produce a full STG which contains the changes, with a minimum amount of time spent making the change.

This led us to take the example of a simple buck converter, and compare what the STG shows, and what the description of operation of the system, in particular it's signals, is. This allowed us to view interactions between certain signals which may not be obvious, but without which the STG would not be a correct representation of the design. For example, Figure 2 shows one scenario of the simple buck converter with some points of interest highlighted.
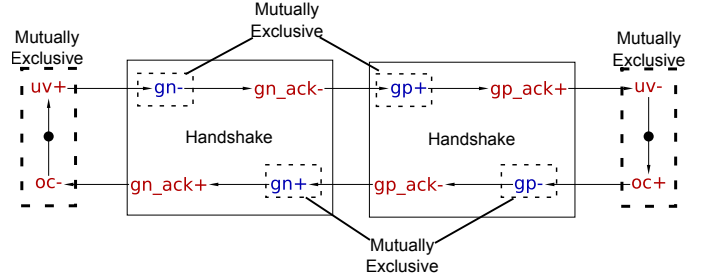

Figure 2: A scenario STG with highlighted points of interest

Studying this STG shows that there are some signal interactions that we can identify. These are:

1) Signals gp and gp_ack form a handshake
2) Signals gn and gn_ack form a handshake
3) Signals uv and oc are mutually exclusive
4) Signals gp and gn are mutually exclusive

These interactions are described in the operation description by the fact that uv and oc indicate opposite conditions in the circuit and will naturally be mutually exclusive, gp and gn are used to switch PMOS and NMOS transistors respectively, and switching these transistors on at the same time will cause a short circuit and should be mutually exclusive, and gp_ack and gn_ack are used to acknowledge the state of these transistors, and thus will follow gp and gn respectively, forming two separate handshakes.

Knowing this information means that we can describe these protocols once, and include them in the design of a circuit involving these signals. Any other interactions between of any of these signals, or any other signals included as part of a circuit will automatically include these protocols, which will avoid circuit breaking bugs in the testing phase. For example, an interaction involving uv will ensure that it is never set high at the same time as oc, as otherwise the mutual exclusion will not hold.

If one of these signals is removed, then the protocols that this affects can be removed, and the unaffected protocols can continue to be used, avoiding a major re-design which may happen when using the monolithic approach.

With this idea in mind, we need to find a way to describe these protocols, as well as other signal interactions, in order to ensure that each different interaction can be edited without needing to change every single interaction.

3. Concepts

In this section we formally introduce *concepts* that we propose to employ for the specification of asynchronous circuits. Below we list (fairly standard) definitions and notational conventions that are used throughout the paper.

We use $\mathbb{B}$ to denote the set of Boolean values $\{0, 1\}$. Given two Boolean functions $f : X \to \mathbb{B}$ and $g : X \to \mathbb{B}$ with the same domain X, we lift Boolean operators (disjunction $\vee$, conjunction $\wedge$, implication $\Rightarrow$, etc.) in the usual manner: $h = f \vee g$ means $h(x) = f(x) \vee g(x)$ for all $x \in X$, etc. Furthermore, $\mathbf{0}$ and $\mathbf{1}$ stand for constant Boolean functions that discard their input and return values 0 and 1, respectively.

A *monoid* is a set M and a binary operation $\diamond : M \times M \to M$ satisfying two axioms:

- Identity: $e \diamond a = a \diamond e = a$ for any $a \in M$, where $e \in M$ is the *identity element* of the monoid.
- Associativity: $a \diamond (b \diamond c) = (a \diamond b) \diamond c$ for all $a, b, c \in M$.

Monoid is the simplest mathematical structure that captures the notions of *emptiness* and *composition*. The concepts introduced in this section form *commutative monoids*: they have identity elements corresponding to empty specifications, and can be composed to build complex concepts from simpler ones. The order of composition does not matter, i.e., the concepts commute: $a \diamond b = b \diamond a$ for all $a, b \in M$.

3.1. Abstract concepts

We first describe *abstract concepts* that we use as building blocks for developing *domain specific concepts*, such as those related to asynchronous circuits (Section 3.2).

Abstract concepts are parameterised by finite sets of *states* S and *events* E. The *initial state concept* captures all possible (or *permitted*) initial states of the system. In the most general form it is a function

$$\mathsf{initial} : S \to \mathbb{B}$$

that given a state $s \in S$ returns 1 if s is an initial state and 0 otherwise. In practice this concept is often realised as a membership test of a set of initial states $I \subseteq S$, i.e. $\mathsf{initial}(s) = s \in I$. However, we prefer the functional form because it is more abstract and permits other, often more efficient realisations. Note that $\mathbf{0}$ and $\mathbf{1}$ have natural interpretations as initial concepts: they correspond to systems with no initial states, and systems where any state can be initial, respectively. Initial state concepts form a commutative monoid with the identity element $\mathbf{1}$ and the composition operation $\wedge$. Intuitively, if a system comprises two subsystems then its initial state should satisfy constraints imposed by both subsystems, hence the conjunction operator.

The *event excitation concept* captures all states wherein a given event can occur (or is *excited*). In the most general form it is a function

$$\mathsf{excited} : E \times S \to \mathbb{B}$$

that given an event $e \in E$ and a state $s \in S$ checks whether e is excited in s. In practice this concept is often realised using

interpreted graph models such as Finite State Machines and Petri Nets [4], Conditional Partial Order Graphs [16], and others. A partial application of the excitation function is often useful: $\mathsf{excited}(e)$ captures all states where event e is excited; for example, if $\mathsf{excited}(e) = \mathbf{0}$ then e is never excited or *dead*. Event excitation concepts also form a commutative monoid with $e = \mathbf{1}$ and $\diamond = \wedge$. This definition corresponds to the *parallel composition* operation, a standard notion for many behavioural models [15].

Some states may be impossible or undesirable during the normal system operation. To express this we use the *invariant concept*, which captures all *correct* or *permitted* states of the system. A typical use case for invariant concepts is to specify assertions or assumptions about the system state space, that may by verified via model checking and/or used for optimising the implementation. In the most general form an invariant concept is a function

$$\mathsf{invariant} : S \to \mathbb{B}$$

that given a state $s \in S$ returns 1 if s is permitted by the invariant and 0 otherwise. Note that if for some state s the initial concept $\mathsf{initial}(s)$ holds but the invariant $\mathsf{invariant}(s)$ does not hold, then the specification is *contradictory* and cannot be satisfied by any implementation. We therefore usually assume that $\mathsf{initial}(s) \Rightarrow \mathsf{invariant}(s)$ holds for all $s \in S$. Similarly, invariant concepts form a commutative monoid with $e = \mathbf{1}$ and $\diamond = \wedge$. Intuitively, if a system comprises two subsystems then its states should be permitted in both of the subsystems.

One can derive other useful concepts from the three concepts described above, for instance,

$$\mathsf{quiescent}(e, s) = \overline{\mathsf{excited}(e, s)}$$

captures all states $s \in S$ when a given event $e \in E$ cannot occur. Furthermore, one can define other useful concepts that cannot be derived from the above, e.g., the *execution concept* capturing the effects that different events have on the system state. Due to space limitations we only consider the three concepts defined above and their derivatives.

All described concepts form monoids, hence their combinations are trivially monoids too. It is therefore convenient to consider triples of concepts $(\mathsf{initial}, \mathsf{excited}, \mathsf{invariant})$ with $(\mathbf{1}, \mathbf{1}, \mathbf{1})$ representing the *empty specification*, and composition $(\mathsf{initial}_1, \mathsf{excited}_1, \mathsf{invariant}_1) \diamond (\mathsf{initial}_2, \mathsf{excited}_2, \mathsf{invariant}_2)$ defined as $(\mathsf{initial}_1 \diamond \mathsf{initial}_2, \mathsf{excited}_1 \diamond \mathsf{excited}_2, \mathsf{invariant}_1 \diamond \mathsf{invariant}_2)$. Importantly, composition of two non-contradictory specifications is always non-contradictory, that is if both $\mathsf{initial}_1(s) \Rightarrow \mathsf{invariant}_1(s)$ and $\mathsf{initial}_2(s) \Rightarrow \mathsf{invariant}_2(s)$ hold for all states $s \in S$, then $\mathsf{initial}_1(s) \diamond \mathsf{initial}_2(s) \Rightarrow \mathsf{invariant}_1(s) \diamond \mathsf{invariant}_2(s)$ holds too.

3.2. Concepts for asynchronous circuits

We now introduce concepts which are specific for the domain of asynchronous circuits and express them using the abstract concepts defined above.

Signal-level concepts: States and events of an asynchronous circuit are parameterised by a fixed set of signals A. A state $s \in S$ is an assignment of Boolean values to signals, i.e. a function $s : A \to \mathbb{B}$, while an event $e \in E$ is a *signal transition*, i.e. a pair $e : A \times \mathbb{B}$ comprising a signal $a \in A$ and the value of the signal *after* the transition occurs. We call transitions $(a, 0)$ and $(a, 1)$ *falling* and *rising*, respectively, and denote them by a^- and a^+ for brevity.

The following two predicates are very useful for constructing concepts:

$$\text{before} : E \times S \to \mathbb{B}$$
$$\text{after} \; : E \times S \to \mathbb{B}$$

A state $s \in S$ is said to be *before* a transition $(a, b) \in E$ if $s(a) \neq b$, i.e. in state s signal a has a value which is different from the resulting value of the transition. Similarly, s is *after* (a, b) if $s(a) = b$ (the transition has already occurred).

We are now ready to define an excitation concept called *consistency* [4]:

$$\text{consistency} = \text{before}$$

This concept captures the requirement that in a consistent asynchronous circuit a signal transition can only be excited in states that are before it.

Another key concept in asynchronous circuits is *causality*: we say that a transition $effect \in E$ causally depends on transition $cause \in E$, denoted as

$$\text{causality}(cause, effect) : E \times S \to \mathbb{B}$$

if $effect$ can occur only in states that are after $cause$. This is an excitation concept, which can be expressed as follows:

$$\text{causality}(cause, effect)(e) = \begin{cases} \mathbf{1} & \text{if } e \neq effect \\ \text{after}(cause) & otherwise \end{cases}$$

In words, we do not add any constraints to events $e \in E$ that are distinct from $effect$, but $effect$ is constrained to occur only after $cause$. Note that function after is used in the partially applied form. We will use a short-hand notation

$$cause \rightsquigarrow effect$$

for the causality concept for convenience.

One can compose two causality concepts using the monoid composition, for example

$$a \rightsquigarrow c \; \diamond \; b \rightsquigarrow c$$

corresponds to so-called AND-causality: event c can only occur after both a and b have occurred. Specifying OR-causality is slightly more tricky:

$$\text{orCausality}(a, b, c)(e) = \begin{cases} \mathbf{1} & \text{if } c \neq e \\ \text{after}(a) \lor \text{after}(b) & otherwise \end{cases}$$

Event c is thus excited after at least one cause has occurred.

Gate-level concepts: Using the causality concept we can express the behaviour of gates in asynchronous circuits. For example, a *buffer* is a gate with one input signal $a \in A$ and one output signal $b \in A$, whose output transitions causally depend on the input ones:

$$\text{buffer}(a, b) = a^+ \rightsquigarrow b^+ \; \diamond \; a^- \rightsquigarrow b^-$$

An *inverter* has a similar conceptual specification, but the output transition is inverted:

$$\text{inverter}(a, b) = a^+ \rightsquigarrow b^- \; \diamond \; a^- \rightsquigarrow b^+$$

A *C-element* is a gate with two inputs a and b and one output c, which synchronises input transitions:

$$\text{cElement}(a, b, c) = a^+ \rightsquigarrow c^+ \diamond b^+ \rightsquigarrow c^+ \diamond a^- \rightsquigarrow c^- \diamond b^- \rightsquigarrow c^-$$

In words, the rising output transition c^+ causally depends on both a^+ and b^+, and the falling output transition c^- causally depends on both a^- and b^-. An alternative way to express the same concept is to reuse the buffer concept:

$$\text{cElement}(a, b, c) = \text{buffer}(a, c) \; \diamond \; \text{buffer}(b, c)$$

Indeed, a C-element combines the constraints imposed on the output transitions by two 'virtual' buffers.

Behaviour of other gates can be similarly defined using concepts, see our Haskell implementation [17].

Protocol-level concepts: In addition to gate-level concepts described above it is often important to specify *protocols* of interaction between multiple gates or components. In this section we demonstrate how one can use concepts to specify asynchronous handshakes and mutual exclusion mechanisms.

Given two signals a and b, a *handshake* between them is the following composition of causality concepts:

$$\text{handshake}(a, b) = a^+ \rightsquigarrow b^+ \diamond b^+ \rightsquigarrow a^- \diamond a^- \rightsquigarrow b^- \diamond b^- \rightsquigarrow a^+$$

Intuitively, we have a two-way asynchronous communication channel, where one party sends transitions a^+ and a^- and the other party responds by corresponding b^+ and b^- transitions. One can notice that the four causality concepts match those found in the buffer and inverter concepts, which leads to an alternative way to express a handshake between a and b:

$$\text{handshake}(a, b) = \text{buffer}(a, b) \diamond \text{inverter}(b, a)$$

Indeed, this conceptual understanding of a handshake as being composed from a buffer and an inverter is often used by circuit designers as a convenient way of reasoning.

In order to specify the initial state of a handshake between signals a and b, we can use functions before and after. For example, $\text{before}(a^+)$ captures the states where signal a is set to 0. We can compose an initial state concept with the handshake concept as follows:

$$\text{handshake00}(a, b) = \text{handshake}(a, b) \diamond \text{before}(a^+) \diamond \text{before}(b^+)$$

The resulting concept corresponds to a handshake between signals a and b that are both initially 0.

The last important concept that requires an introduction

is *mutual exclusion* between two signals a and b:

$$\mathsf{me}(a,b) = a^- \rightsquigarrow b^+ \ \diamond \ b^- \rightsquigarrow a^+ \diamond \overline{\mathsf{after}(a^+) \wedge \mathsf{after}(b^+)}$$

The concept comprises two parts: 1) in terms of causality, we say that rising transitions a^+ and b^+ can only occur after the opposite falling ones, 2) the initial states when $a = b = 1$ are forbidden. Taken together these two parts guarantee that a and b are never set to 1 at the same time, i.e. they are mutually exclusive. We also add $\overline{\mathsf{after}(a^+) \wedge \mathsf{after}(b^+)}$ to the invariant.

We can now specify a *mutual exclusion element* [18] that receives asynchronous requests r_1 and r_2 to a shared resource and grants access to it by corresponding mutually exclusive signals g_1 and g_2:

$$\mathsf{meElement}(r_1, r_2, g_1, g_2) = \mathsf{buffer}(r_1, g_1) \diamond \mathsf{buffer}(r_2, g_2) \diamond \mathsf{me}(g_1, g_2)$$

3.3. Signal types

A discussion of signal types and how these are represented using concepts. The benefits of this, for example we can block input -> input.

4. Circuit specification with concepts

This section presents a method for deriving a circuit specification from a set of concepts that describe its different aspects. We focus on specification of *Speed-Independent* (SI) circuits, which is an important class of asynchronous circuits [6] that work correctly regardless of the gates' delays, while the wires are assumed to have negligible delays. Alternatively, one can regard wire forks as isochronic and add wire delays to the corresponding gate delays (*Quasi-Delay Insensitive* (QDI) circuit class [19]). A convenient formalism for specification of SI circuits is STGs [20], [21], which is a special kind of Petri nets [22] whose transitions are associated with signal events.

4.1. Petri nets and STGs

Formally, a Petri net is defined as a tuple $PN = \langle P, T, F, M_0 \rangle$ comprising finite disjoint sets of *places* P and *transitions* T, *arcs* denoting the flow relation $F \subseteq (P \times T) \cup (T \times P)$ and *initial marking* M_0. There is an arc between $x \in P \cup T$ and $y \in P \cup T$ iff $(x, y) \in F$. The *preset* of a node $x \in P \cup T$ is defined as $\bullet x = \{y \mid (y, x) \in F\}$, and the *postset* as $x \bullet = \{y \mid (x, y) \in F\}$. The dynamic behaviour of a Petri net is defined as a *token game*, changing marking according to the enabling and firing rules. A *marking* is a mapping $M : P \to \mathbb{N}$ denoting the number of *tokens* in each place ($\mathbb{N} = \{0, 1\}$ for *1-safe* Petri nets). A transition t is *enabled* iff $\forall p, p \in \bullet t \Rightarrow M(p) > 0$. The evolution of a Petri net is possible by *firing* the enabled transitions. *Firing* of a transition t results in a new marking M' such that

$$M'(p) = \begin{cases} M(p) - 1 & if \ p \in \bullet t \setminus t\bullet, \\ M(p) + 1 & if \ p \in t \bullet \setminus \bullet t, \\ M(p) & otherwise \end{cases}$$

for all $p \in P$.

An STG is a 1-safe Petri net whose transitions are labelled by signal events, i.e. $STG = \langle P, T, F, M_0, \lambda, Z, v_0 \rangle$, where λ is a *labelling function*, Z is a set of *signals* and $v_0 \in \{0, 1\}^{|Z|}$ is a *vector of initial signal values*. The labelling function $\lambda : T \to Z\pm$ maps transitions into *signal events* $Z\pm = Z \times \{+, -\}$. The signal events labelled $z+$ and $z-$ denote the transitions of signals $z \in Z$ from 0 to 1 (rising edge), or from 1 to 0 (falling edge), respectively. The labelling function does not have to be 1-to-1, i.e. transitions with the same label may occur several times in the net. To distinguish transitions with the same label and refer to them from the text an index $i \in \mathbb{N}$ is attached to their labels as follows: $\lambda(t)/i$, where i differs for different transitions with the same label. STGs inherit the operational semantics of their underlying PNs, including the notions of transition enabling and firing.

Graphically, the places are represented as circles, transitions as text labels, consuming and producing arcs are shown by arrows, and tokens are depicted by dots. For simplicity, the places with one incoming and one outgoing arc are often hidden, allowing arcs (with implicit places) between transitions.

4.2. Composition of concepts

A single concept can be used to describe an initial state, invariant states, a single event or a combination of these, yet describing some protocols using this method can become long winded, as these can involve multiple events. We make use of the monoid composition of concepts to describe complex systems incrementally. Importantly we can mix several levels of system description and refer to signal, gate and protocol level concepts in one specification, depending on which level is more convenient in a particular situation.

Consider a C-element example whose signals a and b are inputs, and signal c is the output. When input a or b changes, we assume it remains in the new state until the output c changes. The following signal-level concepts describe this system:

$$\mathsf{outputRise} = a^+ \rightsquigarrow c^+ \ \diamond \ b^+ \rightsquigarrow c^+$$
$$\mathsf{inputFall} = c^+ \rightsquigarrow a^- \ \diamond \ c^+ \rightsquigarrow b^-$$
$$\mathsf{outputFall} = a^- \rightsquigarrow c^- \ \diamond \ b^- \rightsquigarrow c^-$$
$$\mathsf{inputRise} = c^- \rightsquigarrow a^+ \ \diamond \ c^- \rightsquigarrow b^+$$
$$\mathsf{initialState} = \mathsf{before}(a^+) \diamond \mathsf{before}(b^+) \diamond \mathsf{before}(c^+)$$
$$\mathsf{system} = \mathsf{outputRise} \diamond \mathsf{inputFall} \diamond \mathsf{outputFall} \diamond$$
$$\mathsf{inputRise} \diamond \mathsf{initialState}$$

There are 6 concepts featured, the first 5 of which describe certain operations of the system. The sixth concept composes all of the first 5 concepts, and can be translated to the STG shown in Figure 3. The first four are named according to what they represent, for example, $\mathsf{outputRise}$ describes the events which cause the output to rise. The fifth concept is the initial state concept. This is necessary for the algorithm to produce a scenario STG, as the STG produced will not be usable without knowing the states the signals

will be in when the scenario is entered. The algorithm takes this concept and works out where tokens need to be placed in the STG produced.


Figure 3: STG for the example system.

This set of concepts is only one way of describing this C-element and the environment. Another way could be to use gate-level concepts and describe the environment explicitly. In this case the environment allows the inputs to transition in the opposite direction to the output c, as two inverters would. We can then compose this with the C-element and the same initial state concept:

environment = inverter(c, a) ⋄ inverter(c, b)
system = cElement(a, b, c) ⋄ environment ⋄ initialState

This specification is equivalent to the previous one; indeed one can prove this by rearranging the primitive concepts using the commutativity and associativity axioms of the underlying commutative monoid. Consequently, this specification will be translated to the same STG shown in Figure 3. Figure 4 illustrates all the concepts involved in this specification.

Finally, the designer can also rely on protocol-level concepts, producing the following equivalent specification:

system = handshake00(a, c) ⋄ handshake00(b, c)

This example demonstrates that the presented formal notation for capturing concepts is very flexible and provides the designer with a rich selection of available levels of abstraction, which could be used not only for deriving simplest possible specifications but also for cross-checking the adequacy of specifications by *refactoring* them according to the composition laws.

4.3. Multiple behavioural scenarios

So far we have only considered systems operating in a single behavioural *scenario* specified by a composition of concepts. However, real-life systems often need to support multiple scenarios (e.g., start-up and normal operation, different power modes [8], etc.). This allows each individual scenario to be designed using concepts, and tested individually to ensure they work correctly, before these are combined to produce a full system specification.

To increase the re-usability of scenarios, which helps reduce design time of future systems, this method supports the use of pre-designed scenarios as concepts.

In some cases, a designer may find it easier to split the specification of operational modes further than scenarios and design certain elements separately. In this way, a model may be produced from concepts, which may not be an operational mode on its own, but can be composed and


Figure 4: Example system specified using concepts.

tested separately. In some cases, having several elements predefined using concepts may become useful for quickly designing systems. A predefined logic gate, for example, could be useful to include quickly in any list of concepts when designing multiple scenarios. An STG produced of this element can be referenced in a list of concepts by name (provided that the definition is appropriately imported into the current namespace). When a list of concepts is passed into the STG translation algorithm, all referenced concepts are replaced by the corresponding definitions.

5. Interoperability with STG based tools

Concepts may be a useful way of specifying asynchronous circuits, but specifications need to be verified against certain properties to ensure they will work correctly, and once they are deemed to be correct, they can then be synthesized to produce a circuit implementation. Several standard tools exist which automatically verify and synthesize specifications, such as PETRIFY [4] and MPSAT [5], however these are designed to work with STG specifications. Thus it is necessary to convert concept specifications to STGs in order to be used with these tools.

In this section, we will discuss the algorithm which converts concepts to STGs. Once concepts have been converted, they can be composed and combined using tools, and this will also be discuss this. A method of synthesizing concepts directly also exists, which will be explained, and this gives more options to designers.

5.1. Converting concepts to STGs

INCLUDE A PSEUDOCODE ALGORITHM FOR THIS

As discussed in 4.2, there are multiple ways of representing a specification as concepts. Either using signal-level concepts directly, using gate level concepts, or a combination of concept types. When converting any form of concepts to an STG we pass them to an algorithm which begins by breaking down all concepts into a list of the lowest level of concepts, signal-level concepts. As concepts are compositional, and

$$a^+ \rightsquigarrow c^+$$
$$b^+ \rightsquigarrow c^+$$
$$c^+ \rightsquigarrow a^-$$
$$c^+ \rightsquigarrow b^-$$
$$a^- \rightsquigarrow c^-$$
$$b^- \rightsquigarrow c^-$$
$$c^- \rightsquigarrow a^+$$
$$c^- \rightsquigarrow b^+$$
$$\text{initialise}(a, 0)$$
$$\text{initialise}(b, 0)$$
$$\text{initialise}(c, 0)$$

Figure 5: Concepts for a C-element with environment

gate- and protocol-level concepts are composed of signal-level concepts, this process is simple. Using the example of a C-element with an environment, any representation mentioned in Section 4.2 when passed into this algorithm would produce the list of concepts found in Figure 5.

With this list, we can now start to produce an STG. For this explanation, we will use images to illustrate the example, but the algorithm would produce an STG in the form of a .g file. This file type is compatible with all standard tools, such as PETRIFY, MPSAT and WORKCRAFT.

Initially the algorithm finds all signals in the system, and produces some *consistency loops*. A consistency loop features both the high and low transitions of a signal, and places in between. These loops are used to ensure that an STG satisfies the property of *consistency*, which states that a signal transition can only be excited in states that are before it [4]. The places are used to determine the current state of a signal, and each signal has it's own consistency loop. For the above list of signal-level concepts, the first part of the resulting STG to be produced is shown in Figure 6.


Figure 6: The first step in converting concepts to an STG

Next step in the algorithm is to start connecting the signal transitions in the order of the list of signal-level concepts. First of all, $a^+ \rightsquigarrow c^+$. To represent this in the STG, we need to connect the place after a^+, a_0 in Figure 6, which shows whether a has already transitioned high or not, to the transition of c^+. We use a read-arc to connect these as if a_0 contains a token, a read-arc will allow c^+ to transition

without consuming the token, because this would block $a-$ from occurring.


Figure 7: The second step in converting concepts to STGs

All concepts like this can be added to the STG in the same way. The STG following insertion of all concepts in the above list, before the *initialise* concepts, will produce an STG as shown in Figure 8


Figure 8: The converted STG with all arcs added

At this point, all that is left to add is the initial states. This involves placing tokens in relevant places in the STG. The initial state concepts are $\text{initialise}(a, 0)$, $\text{initialise}(b, 0)$ and $\text{initialise}(c, 0)$. These state that the initial states of all three signals is 0, and therefore we have to add a single token to places a_0, b_0 and c_0. This will produce the fully converted STG, which will look as follows:


Figure 9: The converted C-element with environment STG

The .g file containing this STG can be imported into WORKCRAFT, where it can be viewed, and will look similar to Figure 9, and it can be resynthesized using PETRIFY to produce a more compact and easier to read version of the STG, illustrated in Figure 10

Both Figure 9 and Figure 10 are a correct specification for the C-element with environment. Either of these can be used with WORKCRAFT and existing tools. This allows a designer to visualize their designs as STGs, to test, verify and synthesize, and if there are any corrections or additions to be made these can be done either with the concepts, which can be re-converted to view the updated STG, or to the STG directly.

Figure 10: Resynthesized STG for the example system.

In some cases it may not necessary to visualize a converted STG. A *.g* file is used as input with these tools, without having to import into WORKCRAFT. This can further quicken the process of specification using concepts. This also applies to synthesis, which can be performed using a tool such as PETRIFY or MPSAT, or using *Direct Synthesis*, which is discussed in Section 5.4.

5.2. Parallel composition

With this method of converting concepts to STGs this allows alternative use of tools. While we can specify an entire scenario as concepts and convert this to an STG, it is also possible to convert smaller collections of concepts, or even singular concepts, which are not necessarily a complete STG. These can be visualised, and then composed using *Parallel Composition* within WORKCRAFT using an integrated tool PCOMP [15].

Using the example above, a C-element with environment, we can convert the cElement(e) and the two inverter(e) concepts individually. Including the initial states, this would give us the following STGs.

These STGs can now be passed into PCOMP, which will compose them. The outcome of which will be the same as the concept composition, an STG the same as shown in Figure 10.

5.3. Scenario composition

As mentioned in 4.3, concepts can be used in a composition to produce a scenario, which can be tested and verified separately to allow for easier changes and additions to smaller, less complex, specifications. However, when scenarios are tested and prove to work as required, and satisfy all necessary properties, they then need to be combined to produce a full system specification, so this can be tested and verified, and ultimately be synthesized to produce an implementation.

When combining scenarios, there are several things to consider in how the scenarios fit together. Depending on the application the circuit is being designed for, some scenarios may need to operate in certain orders, for example, one scenario may exist simply to initialise the circuit, therefore this scenario needs to run at start-up, before any of the other scenarios, and then never be run again while the system remains active.

To address this, when combining scenarios we offer some templates, each of which can be used to combine

(a) cElement(a, b, c)

(b) inverter(c, a)


(c) inverter(c, b)

Figure 11: Individual STG conversions of concepts

scenarios in various orders. With this the designer specifies in which ways the scenarios should be combined, and if an order is needed, the order the scenarios should be run from start up. The following are some examples of templates that will be offered before scenarios are combined:

Sequential: Sequential combination will allow a designer to select the order of all scenarios, so when combined, they will run in a sequence. In this case, there may be a clear order in which the scenarios may run, and this needs to be specified by the designer.

Concurrent: In this case there is no order, but one or more of the scenarios in a specification may run in parallel. This template will combine the scenarios in a way that will allow concurrency to occur, limited to requirements of the specification, for instance, the number of scenarios that can be active at any one time, which can be limited by the number of tokens available at once.

Non-deterministic choice: This template will combine the scenarios in a way that allows any of the scenarios to run, but not according to any order to deal with the lack of determinism in the system, using one token which is used only by the running scenario, and this is returned when the scenario completes.

There may be some more complex requirements to combination, and it is possible to combine some scenarios using one template, then including the result in a combination using another template. For example, a system which is non-deterministic may also have a scenario that runs at start

up to initialise the system. In this case, a designer could combine all non-deterministic scenarios first, then combine the resulting STG with the initialising scenario, setting the order so this runs first.

This method of combination can allow for many possible scenario combination styles, and more complex systems can be combined automatically, which in comparison to manual combination, could reduce the number of errors as well as design time.

5.4. Direct synthesis from concepts

It is possible to synthesise directly from STGs, which we call *Direct Synthesis*. Direct synthesis is a method of producing an implementation directly from concepts, finding Boolean functions which are used to produce a logic gate implementation. There are some caveats to this however. Currently, it is suboptimal, and thus the functions it produces may not be the best possible implementation, but it may be a quicker method of finding an implementation. We will explain this algorithm in this section, and give a brief example.

5.4.1. Algorithm of direct synthesis. When finding an implementation for a specification, we need only concern ourselves with the Boolean equations for the output signals. The input signals, the environment, can not be controlled directly by a circuit, it merely reacts to the outputs, but we need to find a way of causing the outputs to changed based on the inputs. In this way we only need to find the Boolean equations for the output signals.

Therefore, we take an output signal, x. With this, we start by finding all concepts in which a signal causes a low to high transition, also known as the set function, x^+, such as $p \rightsquigarrow x^+$, where p is any signal transition. We then record these signals only for x^+ in the form of a Boolean equation, where a low to high signal transition, such as p^+ would be recorded as simply p, and a high to low transition, p^-, would be recorded as $\bar{p}$. All signals which cause this output signal transition in this scenario will be combined using the *AND* ($*$) operation.

This is then repeated to find the equations for the opposite transition for this output signal, in this case the high to low transition, x^-, also known as the reset function. Now we can reduce these equations. First of all, by resolving some Boolean relations. A Boolean relation is when an equation contains several signals which are related, usually by a transition of one being assumed to occur after the transition of another [4]. For example, take a system with two input signals, p and q, and an output signal x. If some concepts in this system are:

$$p^+ \rightsquigarrow x^+$$
$$q^+ \rightsquigarrow x^+$$
$$p^+ \rightsquigarrow q^+$$

The equation for x^+ will contain in some variation $p * q$. These concepts state that both p^+ and q^+ must occur before

x^+ can occur. It also states, however, the timing assumption that p^+ occurs before q^+. It can therefore be assumed that q^+ shows that the p^+ transition has occurred, and can be removed from the Boolean equations.

Further reductions can be performed on these equations using existing algorithms, which will remove any redundancies and simplify them. Following this, the set and reset functions for this signal can then be used for a logic implementation. Depending on the type of synthesis these can be used in different ways, for instance, C-element synthesis uses these set and reset functions as they are, but technology mapping combines these functions by negating only the signals in the reset function, not the equation itself, and then *OR*-ing this with the set function of x^+. This can then be combined with an existing equation for x from previous scenarios by *OR*-ing this with the existing equation.

This can be repeated for each output signal in the scenario, and for each scenario in the specification, and this will produce a full implementation for the system. An algorithm for Direct Synthesis is found in Algorithm 1

Algorithm 1 Algorithm for direct synthesis

---

for Each scenario in system **do**
    for Output signal x in all output signals **do**
        for Each signal transition p that causes x^+ **do**
            if p^+ **then**
                setEq <- setEq $* p$;
            else
                setEq <- setEq $* \bar{p}$
            end if
        end for
        for Each signal transition p that causes x^- **do**
            if p^+ **then**
                resetEq <- resetEq $* p$;
            else
                resetEq <- resetEq $* \bar{p}$
            end if
        end for
    end for
    reduce(setEq)
    reduce(resetEq)
end for

5.4.2. Example of direct synthesis. For this example we will use the same example as we have previously, the C-element with environment. In the case of a circuit for this example, the inputs a and b we cannot control, and thus it is assumed that the environment will correctly invert the output, c to produce these input signals.

Using the list of concepts found in Figure 5, we can find the set and reset functions for the only output signal c.

$$c^+ : a * b$$
$$c^- : \bar{a} * \bar{b}$$

These functions are fairly simple in this example, and there are no further simplification to be done, and thus these

can be used for synthesis. In this case, as we are aiming to use a C- element in the implementation, and thus we can use C-element synthesis. The synthesis of this will produce the implementation found in Figure 12


Figure 12: Implementation synthesized from C-element with environment concepts

6. Case study

In this section we follow a case study to show the design flow of this method using an example from power management domain [3]. A basic power regulator comprises an analogue buck and a digital controller, as shown in Figure 13a. The controller operates the power regulating PMOS and NMOS transistors of the buck (using gp and gn outputs) as a reaction to *under-voltage* (UV), *over-current* (OC) and *zero-crossing* (ZC) conditions (uv, oc and zc inputs, respectively). These conditions are detected and signalled by a set of specialised sensors implemented as comparators of measured current and voltage levels against some reference values (V_ref, I_max, V_0). Note that the gp and gn signals are buffered to drive the very large power regulating transistors and their effect on the buck can be significantly delayed. Therefore, the controller is explicitly notified (by the gp_ack and gn_ack signals) when the power transistor threshold levels (Th_pmos and Th_nmos) are crossed.

The operation of a power regulator is usually specified in an intuitive, but rather informal way, e.g. by enumerating the possible sequences of detected conditions and describing the intended reaction to these events, as shown in Figure 13b. The diagram shows that UV should be handled by switching the NMOS transistor OFF and PMOS transistor ON, while OC should revert their state – PMOS OFF and NMOS ON (no ZC scenario). Detection of the ZC after UV does not change this behaviour (late ZC scenario). However, if ZC is detected before UV then both the PMOS and NMOS transistors remain OFF until the UV condition (early ZC scenario).

Note that ZC and UV are independent conditions that indicate separate physical effects and therefore the corresponding signals can happen in any order. UV indicates that the voltage supplied to the load has decreased below the reference value. ZC occurs when the coil current reduces to 0 and causes the NMOS transistor to switch OFF, so that the NMOS acts like a diode and only conducts in one direction.

6.1. Formal specification

From the informal buck specification and our knowledge of the signals, we can determine three separate operating conditions that occur in the analogue circuit and the controller needs to react to. All of these can be produced from concepts separately and composed to produce scenario

(a) Schematic.
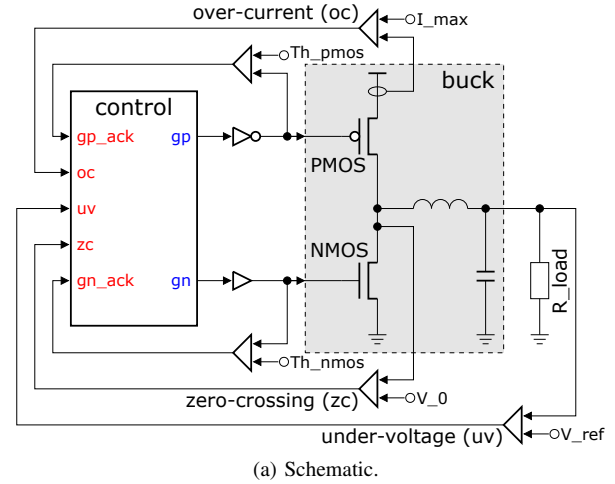

(b) Informal specification.

Figure 13: Buck converter.

STGs, before combining these to produce a single full circuit STG.

During this process, it is useful to find any operations which occur between two or more operational modes, as these can then be reused in other scenarios. If this is defined in one list of concepts, it can then be referenced in concept lists for other scenarios by name.

6.1.1. No ZC scenario. We start with describing the operational mode where no ZC condition is signalled and produce the following list of concepts:

$$\text{uvFunc} = uv^+ \rightsquigarrow gp^+ \diamond uv^+ \rightsquigarrow gn^-$$

$$\text{ocFunc} = oc^+ \rightsquigarrow gp^- \diamond oc^+ \rightsquigarrow gn^+$$

$$\text{uvReact} = gp\_ack^+ \rightsquigarrow uv^- \diamond gn\_ack^+ \rightsquigarrow uv^-$$

$$\text{ocReact} = gp\_ack^- \rightsquigarrow oc^- \diamond gn\_ack^+ \rightsquigarrow oc^-$$

$$\text{environmentConstraint} = \text{me}(uv, oc)$$

$$\text{circuitConstraint} = \text{me}(gn, gp)$$

$$\text{gpHandshake} = \text{handshake}(gp, gp\_ack)$$

$$\text{gnHandshake} = \text{handshake}(gn, gn\_ack)$$

$$\text{initialState} = \text{before}(uv^+) \diamond \text{before}(oc^+)$$

$$\text{chargeFunc} = \text{ocFunc} \diamond \text{ocReact} \diamond$$
$$\text{environmentConstraint} \diamond \text{circuitConstraint} \diamond$$
$$\text{gpHandshake} \diamond \text{gnHandshake} \diamond \text{initialState}$$

$$\text{zcAbsent} = \text{quiescent}(zc^+) \diamond \text{quiescent}(zc^-)$$

$$\text{zcAbsentScenario} = \text{chargeFunc} \diamond \text{uvFunc} \diamond \text{uvReaction} \diamond$$
$$\text{zcAbsent}$$

Figure 14: STG for the zcAbsentScenario concept.

In this list, we introduce concepts which describe the correction of both under-voltage and over-current. This includes concepts describing handshakes for gp/gp_ack and gn/gn_ack, and transistor safety constraints which are described using protocol-level concepts. We also describe constraints provided by the environment that we are aware of, so this behaviour is captured in the resulting STG.

The descriptions of the operational modes suggest that there are similarities between them, mainly in the sequence of PMOS/NMOS activation during the charging cycle. Therefore it is natural to define a chargeFunc concept that can be reused when other operation modes are specified.

Figure 14 shows the STG produced from the final concept in the list, zcAbsentScenario. It contains two tokens which are created based on the initialState concept. The analogue circuit is not guaranteed to signal either under-voltage or over-current first, and this needs to be noted in the concepts and the STG.

6.1.2. Late ZC scenario. In this operational mode we have to include zero-crossing, as per the specification. However, in this case under-voltage and over-current are corrected in the same way as in the no ZC scenario, and therefore we can include the zcAbsentScenario concept. We still need to describe the interactions involving zero-crossing. The concepts are as follows:

$$\text{zcLate} = uv^+ \rightsquigarrow zc^+ \diamond zc^- \rightsquigarrow uv^+$$

$$\text{zcLateScenario} = \text{chargeFunc} \diamond \text{uvFunc} \diamond \text{uvReact} \diamond$$
$$\text{zcLate} \diamond \text{before}(zc^+)$$

Figure 15: STG for the zcLateScenario concept.

Figure 15 shows the STG produced, which looks similar to the STG in Figure 14 but this features a branch from some uv and zc interaction. This is defined in the concept zcLateScenario, which does not describe the arc $zc^+ \rightsquigarrow zc^-$ but this consistency is implied in concepts, as for obvious reasons, zc^+ must occur before zc^- can occur.

6.1.3. Early ZC scenario. The chargeFunc concept can be reused, as the PMOS/NMOS transistors are still operated in the same way. However, the interplay between the early zero-crossing and under-voltage needs to be specified with several new concepts:

$$\text{zcFunc} = zc^+ \rightsquigarrow gn^-$$

$$\text{zcReact} = oc^- \rightsquigarrow zc^+ \diamond gp^+ \rightsquigarrow zc^-$$

$$\text{uvFunc}' = uv^+ \rightsquigarrow gp^+$$

$$\text{uvReact}' = zc^+ \rightsquigarrow uv^+ \diamond zc^- \rightsquigarrow uv^- \diamond gp\_ack^+ \rightsquigarrow uv^-$$

$$\text{zcEarlyScenario} = \text{chargeFunc} \diamond \text{zcFunc} \diamond \text{zcReact} \diamond$$
$$\text{uvFunc}' \diamond \text{uvReact}' \diamond \text{before}(zc^+)$$

Figure 16: STG for the zcEarlyScenario concept.

The obtained STG for the zcEarlyScenario concept is shown in Figure 16. We have now produced three scenario STGs, one for each operational mode in this system. However, we need to ensure they are correct both for the specification and as STG models before we can combine these.

6.2. Verification and simulation of scenarios

To be combinable by our method, the produced STGs need to have certain properties [4]:

- Complete State Coding (CSC): each state of the models with different behaviour has differing signal encodings to avoid problems during synthesis. Note

that in some cases it is possible to automatically resolve a CSC conflict.

- Deadlock freedom: no state is reachable from which no progress can be made.
- Output persistence: there are no race conditions in the STG.
- Signal consistency: in any trace the rising and falling phases of each signal alternate.

These properties are automatically checked in WORKCRAFT using the MPSAT [5] backend tool. In the event that one of these properties does not hold, unless it can be corrected automatically, the composition of scenarios fails. In this case a problematic concept is identified and diagnostic information is printed out to help a designer to correct the issue.

Correctly produced scenarios may not necessarily work as the specification suggests, and this needs to be validated before using these scenarios in any further designs. WORK-CRAFT features a simulation tool, and this can be used by a designer to check that the signals can transition according to the initial requirements. If the simulation produces undesirable results, a designer can work to fix the error of the scenario STG, or correct the design at the concept level. The latter is the preferable method if this design is to be reused either as a predefined concept, or as a scenario in another system.

6.3. Combining scenarios

Now we have scenario STGs that have been verified to ensure they conform to standards required of STGs, and simulated to ensure they work as expected according to the specification. These can now be combined to produce a full system implementation. As mentioned in Section 5.3, we can combine these scenarios in multiple ways, depending on whether there is some sort of required ordering to the way these scenarios must run.

For a simple buck controller there is no required order of running these scenarios, and it is uncertain as to which scenario may be running at any one time. Therefore, the best solution for this system is to combine all of the scenarios in a non-deterministic fashion. The full system implementation produced should allow for only one of these scenarios to run at a time, but when the scenario has completed the system should return to a state where any of the scenarios could run again, regardless of which solution ran previously.

Figure 17 shows the full system specification STG that has been produced from combining the three scenario STGs as seen in Figures 14, 15, and 16. The first most notable part of this full system STG is that there is only one branch for over-current correction. As mentioned above, over-current is corrected in exactly the same way for each of these scenarios, and as such, to reduce the complexity of the model, these can be combined into a single branch that runs after any of the under-voltage correction branches have run.

There are two explicit places in this model, p1 and p2. p1 holds a token initially, and this allows any of the scenarios to run. This place has no control over which scenario can run, but it only allows one of them to run at a
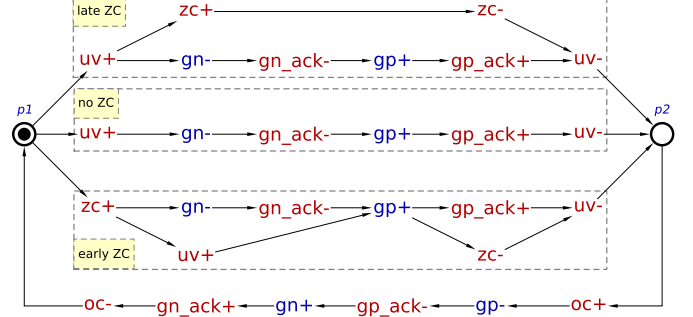

Figure 17: Complete STG for a buck converter.

time. The single token is consumed by whichever scenario runs, and the lack of token in p1 after this stops any more scenarios running. The token is passed through the scenario, and after under-voltage is corrected the token is passed into p2. This token is then consumed by the over-current branch, and returned to p1 and this allows only one of the scenarios to run again.

6.4. Verification and simulation

Like with the scenario models when they have been composed, we need to verify that this model satisfies certain properties after combination of multiple scenarios, as any issues at this stage will cause an implementation of the model to be wrong and this can cause the model to be unimplementable as an SI circuit.

The verification properties we need to satisfy are the same as for scenarios (see Section 6.2), and are corrected in similar ways, however the corrections can be done within the problematic scenario, by changing, adding or removing one or more concepts to avoid affecting any of the functionality of the whole system, or any of the correctly functioning scenarios.

When the full system model satisfies all of the verification properties, we can simulate this model and check that the signals can transition in the order we expect according to the requirements of the system. If this is correct, we can guarantee that this model represents the full system specification, and can now be used in the next step.

6.5. Synthesis of a speed-independent controller

A fully working model of the system is only part way to having completed the process. The final step in this design flow is to synthesize this model. Synthesis is the process of finding Boolean equations to calculate the next state of the output signals based on the input signals and the current state of the circuit [4]. We can do this using PETRIFY or MPSAT, both of which are integrated in WORKCRAFT. Passing this model through one of these tools will produce logic equations that describe how the outputs gp and gn can be produced using uv, oc, zc, gp_ack, gn_ack. Using logic gates, we can reproduce a circuit diagram for these equations. Figure 18 shows the logic circuit, synthesized from this full model.

When we have acquired a circuit design, it needs to be verified to ensure that the logic will perform as we expected

Figure 18: Asynchronous logic gate implementation

before the circuit is fabricated. To verify this, the circuit is converted into a so-called circuit Petri net, using a model for each logic gate and combining them by means of read-arcs. By reachability analysis of this Petri net one can verify that the corresponding circuit is deadlock-free, hazard-free and conforms to the specification [23].

7. Related work

There are several existing methodologies which are similar to the one being proposed in this paper, however we found them limited in certain aspects as discussed below.

Snippets [10], similar to concepts, are smaller state graph models which are used to compose full state graphs of larger systems. Snippets describe the operation of a part of a system in terms of input and output alphabets, and in which ways these snippets can fail. When composed with other snippets it can produce a working system state graph model. With our design methodology however we want to go deeper and decompose a component into concepts responsible for capturing signal behaviours for system features, such as handshakes, mutual exclusion, synchronisation, etc.

DI algebra [12] is a method of describing systems as algebraic equations. Each equation represents an operation of the specification, similar to scenarios, and composing these can be simplified for the most compact version of the equation. These can then be composed to find an equation for the whole specification and again simplified for the most compact version. Our method is similar to DI algebra, however concepts are described textually, which is different to DI algebra and as such, simplification does not occur at concept level, but during the composition and combination steps, and the most compact form of the model is automatically produced. To the best of our knowledge there are no tools or methodologies supporting compositional design of asynchronous circuits using DI algebra and it is therefore not interoperable with the rest of our design flow, and this also makes it unsuitable for use in an industrial setting.

8. Conclusions and future work

This paper shows that it is possible to design a system by splitting it into operational modes, and describing signal interactions and requirements of the mode in a textual format. These can then be used to produce STGs that represent these operational modes, which can be combined to produce a model for the full system specification.

This design method can reduce the time of designing an asynchronous control circuit from the ground up, as well as

allow reuse of components either as part of a scenario or entire scenarios to reduce the design-time of future projects. Composition of concepts and scenarios can help reduce errors and save time in comparison to performing these manually. This method can help to make asynchronous circuits more appealing to industrial designers.

This method currently works with Signal Transition Graphs, however it can be applied to other modelling disciplines, such as Finite State Machines (FSM). In some cases, a designer may wish to view a scenario or full system model as an FSM as they can provide more information about a system, which can help with editing finer details of system, or when correcting errors. It would be possible to use concepts to create scenario FSMs that can be combined to produce a full system model in FSM form. STGs and FSMs could be interchangeable in this respect, for example, a scenario STG could be produced, and a designer may choose to view it as an FSM, which can be viewed by "zooming in" to a section of the STG, that will expand this to show states of the system, and possible transitions. Any edits to this can then be used to update the overall STG.

In addition to FSMs and STGs, we also plan to extend this design method to support Conditional Partial Order Graphs (CPOGs) [16] and Parameterised Graphs [24] for modelling asynchronous circuits. These models are also integrated into WORKCRAFT as part of the SCENCO tool-suite [25], allowing circuits to be described by algebraic equations in text form. SCENCO provides support for describing concept models to produce scenarios and then compose scenarios to produce full system implementations.

In certain cases, an implementation may need to be changed based on design parameters to produce the best result, and these may change during run-time or after the fabrication stage. CPOGs support parameter and run-time reconfigurability [8], hence if design parameters change, the design does not need to be recomposed from another list of concepts, thus saving time. For this reason, CPOGs could work very well with this design approach.

STGs and CPOGs both have their benefits when working with asynchronous systems, and it will be useful to compare these two methods, to see if either is better when designing a system from concepts, or see if there are certain specifications where one modelling method is better than the other.

Acknowledgements

References

[1] Jens Sparsø and Stephen B Furber. *Principles of asynchronous circuit design: a systems perspective.* Springer Netherlands, 2001.

[2] Jonathan Audy. Navigating the path to a successful IC switching regulator design. *Tutorial at IEEE International Solid-State Circuits Conference (ISSCC)*, 2008.

[3] D. Sokolov, A. Mokhov, A. Yakovlev, and D. Lloyd. Towards asynchronous power management. In *IEEE Faible Tension Faible Consommation (FTFC)*, pages 1–4, May 2014.

[4] J. Cortadella, M. Kishinevsky, A. Kondratyev, L. Lavagno, and A. Yakovlev. *Logic Synthesis for Asynchronous Controllers and Interfaces*. Springer, 2002.

[5] Victor Khomenko, Maciej Koutny, and Alex Yakovlev. Detecting state encoding conflicts in stg unfoldings using sat. *Fundamenta Informaticae*, 62(2):221–241, 2004.

[6] W. Bartky D. Muller. A theory of asynchronous circuits. *International Symposium of the Theory of Switching*, 1959.

[7] P. Hudak. Building domain-specific embedded languages. *ACM Computing Surveys*, 28(4):196, 1996.

[8] Andrey Mokhov, Maxim Rykunov, Danil Sokolov, and Alex Yakovlev. Design of processors with reconfigurable microarchitecture. *Journal of Low Power Electronics and Applications*, 4(1):26–43, 2014.

[9] I. Poliakov, D. Sokolov, and A. Mokhov. Workcraft: a static data flow structure editing, visualisation and analysis tool. In *Petri Nets and Other Models of Concurrency*, pages 505–514. 2007.

[10] Igor Benko and Jo Ebergen. Composing snippets. In Jordi Cortadella, Alex Yakovlev, and Grzegorz Rozenberg, editors, *Concurrency and Hardware Design*, volume 2549 of *Lecture Notes in Computer Science*, pages 1–33. Springer Berlin Heidelberg, 2002.

[11] Craig Armenti. Get to market faster with modular circuit design. *Electronic Engineering Journal*, 2015.

[12] Mark B Josephs and Jan Tijmen Udding. An overview of di algebra. In *Hawaii International Conference on System Sciences (HICSS)*, volume 1, pages 329–338. IEEE, 1993.

[13] Arseniy Alekseyev, Ivan Poliakov, Victor Khomenko, and Alex Yakovlev. Optimisation of balsa control path using stg resynthesis.

[14] Luis A Plana, Sam Taylor, and Doug Edwards. Attacking control overhead to improve synthesised asynchronous circuit performance. In *IEEE International Conference on Computer Design (ICCD)*, pages 703–710, 2005.

[15] Arseniy Alekseyev, Victor Khomenko, Andrey Mokhov, Dominic Wist, and Alex Yakovlev. Improved parallel composition of labelled petri nets. In *International Conference on Application of Concurrency to System Design (ACSD)*, pages 131–140, 2011.

[16] Andrey Mokhov and Alex Yakovlev. Conditional partial order graphs: Model, synthesis, and application. *IEEE Transactions on Computers*, 59(11):1480–1493, 2010.

[17] Concepts repository. https://github.com/snowleopard/concepts, 2015.

[18] D. J. Kinniment. *Synchronization and Arbitration in Digital Systems*. John Wiley and Sons, 2008. ISBN: 978-0-470-51082-7.

[19] A. Martin. Compiling communicating processes into delay-insensitive vlsi circuits. *Distributed Computing, vol. 1(4)*, pages 226–234, 1986.

[20] T.-A. Chu. *Synthesis of self-timed VLSI circuits from graph-theoretic specifications*. PhD thesis, Massachusetts Institute of Technology, 1987.

[21] A. Yakovlev L. Rosenblum. Signal graphs: from self-timed to timed ones. *International Workshop on Timed Petri Nets*, pages 199–206.

[22] C. Petri. *Kommunikation mit automaten (Communicating with automata)*. PhD thesis, University of Bonn, 1962.

[23] Ivan Poliakov, Andrey Mokhov, Ashur Rafiev, Danil Sokolov, and Alex Yakovlev. Automated verification of asynchronous circuits using circuit Petri nets. In *IEEE International Symposium on Asynchronous Circuits and Systems (ASYNC)*, pages 161–170, 2008.

[24] Andrey Mokhov and Victor Khomenko. Algebra of parameterised graphs. *ACM Transactions on Embedded Computing Systems*, 13(4s), 2014.

[25] SCENCO toolsuite. http://www.workcraft.org/scenco, 2015.